



lavaan: An R Package for Structural Equation Modeling

Yves Rosseel
Ghent University

Abstract

Structural equation modeling (SEM) is a vast field and widely used by many applied researchers in the social and behavioral sciences. Over the years, many software packages for structural equation modeling have been developed, both free and commercial. However, perhaps the best state-of-the-art software packages in this field are still closed-source and/or commercial. The R package **lavaan** has been developed to provide applied researchers, teachers, and statisticians, a free, fully open-source, but commercial-quality package for latent variable modeling. This paper explains the aims behind the development of the package, gives an overview of its most important features, and provides some examples to illustrate how **lavaan** works in practice.

Keywords: structural equation modeling, latent variables, path analysis, factor analysis.

1. Introduction

This paper describes package **lavaan**, a package for structural equation modeling implemented in the R system for statistical computing (R Development Core Team 2012). The package is available from the Comprehensive R Archive Network (CRAN) at <http://CRAN.R-project.org/package=lavaan> and supported by the website <http://lavaan.org/>. **lavaan** is an acronym for *latent variable analysis*, and its name reveals the long-term goal: to provide a collection of tools that can be used to explore, estimate, and understand a wide family of latent variable models, including factor analysis, structural equation, longitudinal, multilevel, latent class, item response, and missing data models (Skrondal and Rabe-Hesketh 2004; Lee 2007; Muthén 2002).

However, the development of **lavaan** has only begun and much remains to be done to reach this ambitious goal. To date, the development of **lavaan** has focused on structural equation modeling (SEM) with continuous observed variables (Bollen 1989), which is the focus of this

paper. Structural equation models encompass a wide range of multivariate statistical techniques. The history of the field traces back to three different traditions: (1) path analysis, originally developed by the geneticist Sewall Wright (Wright 1921), later picked up in sociology (Duncan 1966), (2) simultaneous-equation models, as developed in economics (Haavelmo 1943; Koopmans 1945), and (3) factor analysis from psychology (Spearman 1904; Lawley 1940; Anderson and Rubin 1956). The three traditions were ultimately merged in the early 1970s and although many different researchers have made significant contributions (Jöreskog 1970; Hauser and Goldberger 1971; Zellner 1970; Keesling 1973; Wiley 1973; Browne 1974), it was the work of Karl Jöreskog (Jöreskog 1973), that came to dominate the field. Not least because he (together with Dag Sörbom) developed a computer program called **LISREL** (for LInear Structural RELations), providing many applied researchers access to this new and exciting field of structural equation modeling. From 1974 onwards, **LISREL** was distributed commercially by Scientific Software International. In the following decades, the wide availability of **LISREL** initiated a methodological revolution in the social and behavioral sciences. Today, almost four decades later, **LISREL** 8 (Jöreskog and Sörbom 1997) is still one of the most widely used software packages for structural equation modeling.

In the years after the birth of **LISREL**, many technical advances were made and several new software packages for structural equation modeling were developed. Some of the more popular ones that are still in wide use today are **EQS** (Bentler 2004), **AMOS** (Arbuckle 2011) and **Mplus** (Muthén and Muthén 2010), all of which are commercial. The few non-commercial SEM programs outside the R environment are **Mx** (Neale, Boker, Xie, and Maes 2003) (free, but closed-source), and **gllamm**, which is implemented in *Stata* (Rabe-Hesketh, Skrondal, and Pickles 2004).

Within the R environment, there are two approaches to estimate structural equation models. The first approach is to connect R with external commercial SEM programs. This is often useful in simulation studies where fitting a model with SEM software is one part of the simulation pipeline (see, for example, van de Schoot, Hoijtink, and Deković 2010). During one run of the simulation, syntax is written to a file in a format that can be read by the external SEM program (say, **Mplus** or **EQS**); the model is fitted by the external SEM program and the resulting output is written to a file that needs to be parsed manually to extract the relevant information for the study at hand. Depending on the SEM program, the connection protocols can be tedious to set up. Fortunately, two R packages have been developed to ease this process: **MplusAutomation** (Hallquist 2012) and **REQS** (Mair, Wu, and Bentler 2010), to communicate with the **Mplus** and **EQS** program respectively. The second approach is to use a dedicated R package for structural equation modeling. At the time of writing, apart from **lavaan**, there are two alternative packages available. The **sem** package, developed by John Fox, has been around since 2001 (Fox, Nie, and Byrnes 2012; Fox 2006) and for a long time, it was the only package for SEM in the R environment. The second package is **OpenMx** (Boker, Neale, Maes, Wilde, Spiegel, Brick, Spies, Estabrook, Kenny, Bates, Mehta, and Fox 2011), available from <http://openmx.psyc.virginia.edu/>. As the name of the package suggests, **OpenMx** is a complete rewrite of the **Mx** program, consisting of three parts: a front end in R, a back end written in C, and a third-party commercial optimizer (NPSOL). All parts of **OpenMx** are open-source, except of course the NPSOL optimizer, which is closed-source.

The rest of the paper is organized as follows. First, I describe why I began developing **lavaan** and how my initial objectives impacted the software design. Next, I illustrate the most characteristic feature of **lavaan**: the ‘**lavaan** model syntax’. In the sections that follow,

I present two well-known examples from the SEM literature (a CFA example, and a SEM example) to illustrate the use of **lavaan** in practice. Next, I discuss the use of multiple groups, and in the last section before the conclusion, I provide a brief summary of features included in **lavaan** that may be of interest to applied researchers.

2. Why do we need lavaan?

As described above, many SEM software packages are available, both free and commercial, including a couple of packages that run in the R environment. Why then is there a need for yet another SEM package? The answers to this question are threefold:

1. **lavaan** aims to appeal to a large group of applied researchers that needs SEM software to answer their substantive questions. Many applied researchers have not previously used R and are accustomed to commercial SEM programs. Applied researchers often value software that is intuitive and rich with modeling features, and **lavaan** tries to fulfill both of these objectives.
2. **lavaan** aims to appeal to those who teach SEM classes or SEM workshops; ideally, teachers should have access to an easy-to-use, but complete, SEM program that is inexpensive to install in a computer classroom.
3. **lavaan** aims to appeal to statisticians working in the field of SEM. To implement a new methodological idea, it is advantageous to have access to an open-source SEM program that enables direct access to the SEM code.

The first aim is arguably the most difficult one to achieve. If we wish to convince users of commercial SEM programs to use **lavaan**, there must be compelling reasons to switch. That **lavaan** is free is often irrelevant. What matters most to many applied researchers is that (1) the software is easy and intuitive to use, (2) the software has all the features they want, and (3) the results of **lavaan** are very close, if not identical, to those reported by their current commercial program. To ensure that the software is easy and intuitive to use, I developed the ‘**lavaan** model syntax’ which provides a concise approach to fitting structural equation models. Two features that many applied researchers often request are support for non-normal (but continuous) data, and handling of missing data. Both features have received careful attention in **lavaan**. And lastly, to ensure that the results reported by **lavaan** are comparable to the output of commercial programs, all fitting functions in **lavaan** contain a `mimic` option. If `mimic = "Mplus"`, **lavaan** makes an effort to produce output that resembles the output of **Mplus**, both numerically and visually. If `mimic = "EQS"`, **lavaan** produces output that approaches the output of **EQS**, at least numerically (not visually). In future releases of **lavaan**, we plan to add `mimic = "LISREL"` and `mimic = "AMOS"` (but users of those programs can currently use `mimic = "EQS"` as a proxy for those).

The second aim targets those of us that teach SEM techniques in classes or workshops. For teachers, the fact that **lavaan** is free *is* important. If the software is free, there is no longer a need to install limited ‘student-versions’ of the commercial programs to accompany the SEM course. Of course, teachers will also appreciate an easy and intuitive user experience, so that they can spend more time discussing and interpreting the substantive results of a SEM analysis, instead of expending time explaining the awkward model syntax of a specific program.

Finally, the `mimic` option makes a smooth transition possible from **lavaan** to one of the major commercial programs, and back. Therefore, students who received initial instruction in SEM with **lavaan** should have little difficulty using other (commercial) SEM programs in the future.

The third aim targets professional statisticians working in the field of structural equation modeling. For too long, this field has been dominated by closed-source commercial software. In practice, this meant that many of the technical contributions in the field were realized by those research groups (and their collaborators) that had access to the source code of one of the commercial programs. They could use the infrastructure that was already there to implement and evaluate their newest ideas. Outsiders were forced to write their own software. Some of them, faced with the enormous time-investment that is needed for writing SEM software from scratch, may have given up, and changed their research objectives altogether. Indeed, it seems unfortunate that new developments in this field have potentially been hindered by the lack of open-source software that researchers can use, and reuse, to bring computational and statistical advances to the field. This is in sharp contrast to other fields such as statistical genetics or neuroimaging, where nearly exponential progress has been made in part because both fields rely heavily on, and are driven by, open-source packages. Therefore, I chose to keep **lavaan** fully open-source, without any dependencies on commercial and/or closed-source components. In addition, the design of **lavaan** is extremely modular. Adding a new function for computing standard errors, for example, would require just two steps: (1) adding the new function to the source file containing all the other functions for computing various types of standard errors, and (2) adding an option to the `se` argument in the fitting functions of **lavaan**, allowing the user to request this new type of standard errors.

3. From model to syntax

Path diagrams are often a starting point for applied researchers seeking to fit a SEM model (see Figure 2 for an example). Informally, a path diagram is a schematic drawing that represents a concise overview of the model the researcher aims to fit. It includes all the relevant observed variables (typically represented by square boxes) and the latent variables (represented by circles), with arrows that illustrate the (hypothesized) relationships among these variables. A direct effect of one variable on another is represented by a single-headed arrow, while (unexplained) correlations between variables are represented by double-headed arrows. The main problem for the applied researcher is typically to convert this diagram into the appropriate input expected by the SEM program. In addition, the researcher has to take extra care to ensure the model is identifiable and estimable.

3.1. Specifying models in commercial SEM programs

In the early days of SEM, the only way to specify a model was by setting up the model matrices directly. This was the case for **LISREL**, and many generations of SEM users (including the author of this paper) have come to associate the practice of SEM modeling with setting up a **LISREL** syntax file. This required a good grasp of the underlying theory, and – for some – an incentive to review the Greek alphabet once more. For many first time users, the translation of their diagram directly to **LISREL** syntax was an unpleasant experience. And it added to the still wide-spread belief that SEM modeling is something that should be left to experts, well-versed in matrix algebra (and the Greek alphabet).

In the mid-1980s, **EQS** was the first program to offer a matrix-free model specification. The **EQS** model syntax distinguishes among four fundamental variable types: (1) measured variables, (2) latent variables or factors, (3) measured variable residuals or errors, and (4) latent variable residuals or disturbances. The four types are labeled V, F, E and D respectively. Rather than providing a full model matrix specification, users needed only to identify these four types of variables and their relations. For many applied researchers, this was a giant leap forward, and the **EQS** program quickly became successful. Soon after, this regression-oriented approach was adopted by many other programs (including **LISREL**, which introduced the SIMPLIS language with **LISREL** 8).

In the 1990s, the rise of operating systems with a graphical user interface led to a new evolution in the SEM world. The **AMOS** program, originally developed by James L. Arbuckle, offered a comprehensive graphical interface that allowed users to specify their model by drawing its path diagram. There is no doubt that this approach was very appealing to many SEM users, and again, many commercial SEM packages (including **EQS** and **LISREL**) added similar capabilities to their programs.

But a pure graphical approach is not without its limitations. Sometimes, it can be very tedious to draw each and every element of a path diagram, especially for large models. In addition, many (advanced) features do not translate easily in a graphical environment. For example, how do you specify nonlinear inequality constraints without relying on additional syntax? Although a graphical interface may be excellent as a teaching tool, or as an entry point for first-time users, an accessible text-based syntax may ultimately be more convenient. This is the approach used by **Mplus**. In the **Mplus** program, no graphical interface is available to specify the model, yet many models can be specified in a very concise and compact way. Only the core measurement and structural parameters of a model need to be specified. For example, in **Mplus**, there is no need to list all the residual variances that are part of the model. **Mplus** will add these parameters automatically, keeping the syntax short and easy to understand.

3.2. Specifying models in lavaan

In the **lavaan** package, models are specified by means of a powerful, easy-to-use text-based syntax describing the model, referred to as the ‘lavaan model syntax’. Consider a simple regression model with a continuous dependent variable y , and four independent variables x_1 , x_2 , x_3 and x_4 . The usual regression model can be written as follows:

$$y_i = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + \beta_3 x_{3i} + \beta_4 x_{4i} + \epsilon_i$$

where β_0 is called the intercept, β_1 to β_4 are the regression coefficients for each of the four variables, and ϵ_i is the residual error for observation i . One of the attractive features of the R environment is the compact way we can express a regression formula like the one above:

```
y ~ x1 + x2 + x3 + x4
```

In this formula, the tilde sign (`~`) is the regression operator. On the left-hand side of the operator, we have the dependent variable (y), and on the right-hand side, we have the independent variables, separated by the `+` operator. Note that the intercept is not explicitly included in the formula. Nor is the residual error term. But when this model is fitted (say,

using the `lm()` function), both the intercept and the variance of the residual error will be estimated. The underlying logic, of course, is that an intercept and residual error term are (almost) always part of a (linear) regression model, and there is no need to mention them in the regression formula. Only the structural part (the dependent variable, and the independent variables) needs to be specified, and the `lm()` function takes care of the rest.

One way to look at SEM models is that they are simply an extension of linear regression. A first extension is that you can have several regression equations at the same time. A second extension is that a variable that is an independent (exogenous) variable in one equation can be a dependent (endogenous) variable in another equation. It seems natural to specify these regression equations using the same syntax as used for a single equation in R; we only have more than one of them. For example, we could have a set of three regression equations:

```
y1 ~ x1 + x2 + x3 + x4
y2 ~ x5 + x6 + x7 + x8
y3 ~ y1 + y2
```

This is the approach taken by **lavaan**. Multiple regression equations are simply a set of regression formulas, using the typical syntax of an R formula.

A third extension of SEM models is that they include continuous latent variables. In **lavaan**, any regression formula can contain latent variables, both as a dependent or as an independent variable. For example, in the syntax shown below, the variables starting with an ‘f’ are latent variables:

```
y ~ f1 + f2 + x1 + x2
f1 ~ x1 + x2
```

This part of the model syntax would correspond with the ‘structural part’ of a SEM model. To describe the ‘measurement part’ of the model, we need to specify the (observed or latent) indicators for each of the latent variables. In **lavaan**, this is done with the special operator ‘`=~`’, which can be read as *is manifested by*. The left-hand side of this formula contains the name of the latent variable. The right-hand side contains the indicators of this latent variable, separated by the ‘`+`’ operator. For example:

```
f1 =~ item1 + item2 + item3
f2 =~ item4 + item5 + item6 + item7
f3 =~ f1 + f2
```

In this example, the variables `item1` to `item7` are observed variables. Therefore, the latent variables `f1` and `f2` are first-order factors. The latent variable `f3` is a second-order factor, since all of its indicators are latent variables themselves.

To specify (residual) variances and covariances in the model syntax, **lavaan** provides the ‘`~~`’ operator. If the variable name at the left-hand side and the right-hand side are the same, it is a variance. If the names differ, it is a covariance. The distinction between residual (co)variances and non-residual (co)variances is made automatically. For example:

```
item1 ~~ item1    # variance
item1 ~~ item2    # covariance
```

Formula type	Operator	Mnemonic
Latent variable	=~	is manifested by
Regression	~	is regressed on
(Residual) (co)variance	~~	is correlated with
Intercept	~ 1	intercept
Defined parameter	:=	is defined as
Equality constraint	==	is equal to
Inequality constraint	<	is smaller than
Inequality constraint	>	is larger than

Table 1: Top panel of the table contains the four formula types that can be used to specify a model in the **lavaan** model syntax. The lower panel contains additional operators that are allowed in the **lavaan** model syntax.

Finally, intercepts for observed and latent variables are simple regression formulas (using the ‘~’ operator) with only an intercept (explicitly denoted by the number ‘1’) as the only predictor:

```
item1 ~ 1    # intercept of an observed variable
f1 ~ 1      # intercept of a latent variable
```

Using these four *formula types*, a large variety of latent variable models can be described. For reference, we summarize the four formula types in the top panel of Table 1.

A typical model syntax describing a SEM model will contain multiple formula types. In **lavaan**, to glue them together, they must be specified as a literal string. In the R environment, this can be done by enclosing the formula expressions with (single) quotes. For example,

```
myModel <- '# regressions
  y ~ f1 + f2
  y ~ x1 + x2
  f1 ~ x1 + x2
# latent variables
  f1 =~ item1 + item2 + item3
  f2 =~ item4 + item5 +
    item6 + item7
  f3 =~ f1 + f2

# (residual) variances and covariances
  item1 ~~ item1
  item1 ~~ item2
# intercepts
  item1 ~ 1
  f1 ~ 1'
```

This piece of code will produce a model syntax object called `myModel` that can be used later when calling a function that estimates this model given a dataset, and it illustrates several

features of the **lavaan** model syntax. Formulas can be split over multiple lines, and you can use comments (starting with the `#` character) and blank lines within the single quotes to improve readability of the model syntax. The order in which the formulas are specified does not matter. Therefore, you can use the latent variables in the regression formulas even before they are defined by the `'=~'` operator. And finally, since this model syntax is nothing more than a literal string, you can type the syntax in a separate text file and use a function like `readLines()` to read it in. Alternatively, the text processing infrastructure of R may be used to generate the syntax for a variety of models, perhaps when running a large simulation study.

4. A first example: Confirmatory factor analysis

The **lavaan** package contains a built-in dataset called `HolzingerSwineford1939`. We therefore start with loading the **lavaan** package:

```
R> library("lavaan")
```

The Holzinger & Swineford 1939 dataset is a ‘classic’ dataset that has been used in many papers and books on structural equation modeling, including some manuals of commercial SEM software packages. The data consists of mental ability test scores of seventh- and eighth-grade children from two different schools (Pasteur and Grant-White). In our version of the dataset, only 9 out of the original 26 tests are included. A CFA model that is often proposed for these 9 variables consists of three correlated latent variables (or factors), each with three indicators:

- a *visual* factor measured by 3 variables: `x1`, `x2` and `x3`,
- a *textual* factor measured by 3 variables: `x4`, `x5` and `x6`,
- a *speed* factor measured by 3 variables: `x7`, `x8` and `x9`.




Figure 1: Path diagram of the three factor model for the Holzinger & Swineford data.

id	lhs	op	rhs	user	free	ustart
1	visual	=~	x1	1	0	1
2	visual	=~	x2	1	1	NA
3	visual	=~	x3	1	2	NA
4	textual	=~	x4	1	0	1
5	textual	=~	x5	1	3	NA
6	textual	=~	x6	1	4	NA
7	speed	=~	x7	1	0	1
8	speed	=~	x8	1	5	NA
9	speed	=~	x9	1	6	NA
10	x1	~~	x1	0	7	NA
11	x2	~~	x2	0	8	NA
12	x3	~~	x3	0	9	NA
13	x4	~~	x4	0	10	NA
14	x5	~~	x5	0	11	NA
15	x6	~~	x6	0	12	NA
16	x7	~~	x7	0	13	NA
17	x8	~~	x8	0	14	NA
18	x9	~~	x9	0	15	NA
19	visual	~~	visual	0	16	NA
20	textual	~~	textual	0	17	NA
21	speed	~~	speed	0	18	NA
22	visual	~~	textual	0	19	NA
23	visual	~~	speed	0	20	NA
24	textual	~~	speed	0	21	NA

Table 2: A complete list of all parameters in the three-factor CFA model for the Holzinger & Swineford data.

In what follows, we will refer to this 3 factor model as the ‘H&S model’, graphically represented in Figure 1. Note that the path diagram in the figure is simplified: it does not indicate the residual variances of the observed variables or the variances of the exogenous latent variables. Still, it captures the essence of the model. Before discussing the **lavaan** model syntax for this model, it is worthwhile first to identify the free parameters in this model. There are three latent variables (factors) in this model, each with three indicators, resulting in nine factor loadings that need to be estimated. There are also three covariances among the latent variables – another three parameters. These 12 parameters are represented in the path diagram as single-headed and double-headed arrows, respectively. In addition, however, we need to estimate the residual variances of the nine observed variables and the variances of the latent variables, resulting in 12 additional free parameters. In total we have 24 parameters. But the model is not yet identified because we need to set the metric of the latent variables. There are typically two ways to do this: (1) for each latent variable, fix the factor loading of one of the indicators (typically the first) to a constant (conventionally, 1.0), or (2) standardize the variances of the latent variables. Either way, we fix three of these parameters, and 21 parameters remain free. Table 2, produced by the **parTable()** method, contains an overview of all the relevant parameters for this model, including three fixed factor

loadings. Each row in the table corresponds to a single parameter. The ‘**rhs**’, ‘**op**’ and ‘**lhs**’ columns uniquely define the parameters of the model. All parameters with the ‘**=~**’ operator are factor loadings, whereas all parameters with the ‘**~~**’ operator are variances or covariances. The nonzero elements in the ‘**free**’ column are the free parameters of the model. The zero elements in the ‘**free**’ column correspond to fixed parameters, whose value is found in the ‘**ustart**’ column. The meaning of the ‘**user**’ column will be explained below.

4.1. Specifying a model using the **lavaan** model syntax

There are three approaches in **lavaan** to specify a model. In the first approach, a minimal description of the model is given by the user and the remaining elements are added automatically by the program. This ‘user-friendly’ approach is implemented in the fitting functions **cfa()** and **sem()**. In the second approach, a complete explication of all model parameters must be provided by the user – nothing is added automatically. This is the ‘power-user’ approach, implemented in the function **lavaan()**. Finally, in a third approach, the minimalist and complete approaches are blended by providing an incomplete description of the model in the model syntax, but adding selected groups of parameters using the **auto.*** arguments of the **lavaan** function. We illustrate and discuss each of these approaches in turn.

*Method 1: Using the **cfa()** and **sem()** functions*

In the first approach, the idea is that the model syntax provided by the user should be as concise and intelligible as possible. To accomplish this, typically only the latent variables (using the ‘**=~**’ operator) and regressions (using the ‘**~**’ operator) are included in the model syntax. The other model parameters (for this model: the residual variances of the observed variables, the variances of the factors and the covariances among the factors) are added automatically. Since the H&S example contains three latent variables, but no regressions, the minimalist syntax is very short:

```
R> HS.model <- 'visual  =~ x1 + x2 + x3
+                  textual =~ x4 + x5 + x6
+                  speed   =~ x7 + x8 + x9'
```

We can now fit the model as follows:

```
R> fit <- cfa(HS.model, data = HolzingerSwineford1939)
```

The function **cfa()** is a dedicated function for fitting confirmatory factor analysis (CFA) models. The first argument is the object containing the **lavaan** model syntax. The second argument is the dataset that contains the observed variables. The ‘**user**’ column in Table 2 shows which parameters were explicitly contained in the user-specified model syntax (= 1), and which parameters were added by the **cfa()** function (= 0). If a model has been fitted, it is always possible (and highly informative) to inspect this parameter table by using the following command:

```
parTable(fit)
```

When using the **cfa()** (or **sem()**) function, several sets of parameters are included by default. A complete list of these parameter sets is provided in the top panel of Table 3. In

Keyword	Operator	Parameter set
<code>auto.var</code>	<code>~~</code>	(residual) variances observed and latent variables
<code>auto.cov.y</code>	<code>~~</code>	(residual) covariances observed and latent endogenous variables
<code>auto.cov.lv.x</code>	<code>~~</code>	covariances among exogenous latent variables
Keyword	Default	Action
<code>auto.fix.first</code>	<code>TRUE</code>	fix the factor loading of the first indicator to 1
<code>auto.fix.single</code>	<code>TRUE</code>	fix the residual variance of a single indicator to 0
<code>int.ov.free</code>	<code>TRUE</code>	freely estimate the intercepts of the observed variables (only if a mean structure is included)
<code>int.lv.free</code>	<code>FALSE</code>	freely estimate the intercepts of the latent variables (only if a mean structure is included)

Table 3: Top panel: Sets of parameters that are automatically added to the model representation when the functions `cfa()` or `sem()` are used. Bottom panel: The set of actions automatically taken in an attempt to fulfill the minimum requirements for an identifiable model. These defaults are used by the `cfa()` or `sem()` functions only.

addition, several steps are taken in an attempt to fulfill the minimum requirements for an identifiable model. These steps are listed in the bottom panel of Table 3. In our example, only the first action (fixing the factor loadings of the first indicator) is used. The second one (`auto.fix.single`) is only needed if the model contains a latent variable that is manifested by a single indicator. The third and the fourth actions (`int.ov.free` and `int.lv.free`, respectively) are only needed if a mean structure is added to the model.

Before we move on to the next method, it is important to stress that all of these ‘automatic’ actions can be overridden. The model syntax always has precedence over the automatically generated actions. If, for example, one wishes not to fix the factor loadings of the first indicator, but instead to fix the variances of the latent variances, the model syntax would be adapted as follows:

```
R> HS.model.bis <- 'visual  =~ NA*x1 + x2 + x3
+                      textual =~ NA*x4 + x5 + x6
+                      speed   =~ NA*x7 + x8 + x9
+                      visual  ~~ 1*visual
+                      textual  ~~ 1*textual
+                      speed    ~~ 1*speed'
```

As illustrated above, model parameters are fixed by pre-multiplying them with a numeric value, and otherwise fixed parameters are freed by pre-multiplying them with ‘NA’. The model syntax above overrides the default behavior of fixing the first factor loading and estimating the factor variances. In practice, however, a much more convenient method to use this parameterization is to keep the original syntax, but add the `std.lv = TRUE` argument to the `cfa()` function call:

```
R> fit <- cfa(HS.model, data = HolzingerSwineford1939, std.lv = TRUE)
```

Method 2: Using the lavaan() function

In many situations, using the concise model syntax in combination with the `cfa()` and `sem()` functions is extremely convenient, particularly for many conventional models. But sometimes, these automatic actions may get in the way, especially when non-standard models need to be specified. For these situations, users may prefer to use the `lavaan()` function instead. The `lavaan()` function has the ‘feature’ that it does not add any extra parameters to the model by default, nor does it attempt to make the model identifiable. If the `lavaan()` function is called without any use of the `auto.*` arguments, it becomes the user’s responsibility to specify the correct model syntax. This can lead to lengthier model specifications, but the user has full control. For the H&S model, the full **lavaan** model syntax would be:

```
R> HS.model.full <- '# latent variables
+               visual  =~ 1*x1 + x2 + x3
+               textual =~ 1*x4 + x5 + x6
+               speed   =~ 1*x7 + x8 + x9
+               # residual variances observed variables
+               x1 ~~ x1
+               x2 ~~ x2
+               x3 ~~ x3
+               x4 ~~ x4
+               x5 ~~ x5
+               x6 ~~ x6
+               x7 ~~ x7
+               x8 ~~ x8
+               x9 ~~ x9
+               # factor variances
+               visual  ~~ visual
+               textual  ~~ textual
+               speed    ~~ speed
+               # factor covariances
+               visual  ~~ textual + speed
+               textual  ~~ speed'
R> fit <- lavaan(HS.model.full, data = HolzingerSwineford1939)
```

Method 3: Using the lavaan() function with the auto. arguments*

When using the `lavaan()` function, the user has full control, but the model syntax may become long and contain many formulas that could easily be added automatically. To compromise between a complete model specification using **lavaan** syntax and the automatic addition of certain parameters, the `lavaan()` function provides several optional arguments that can be used to add a particular set of parameters to the model, or to fix a particular set of parameters (see Table 3). For example, in the model syntax below, the first factor loadings are explicitly fixed to one, and the covariances among the factors are added manually. It would be more convenient and concise, however, to omit the residual variances and factor variances from the model syntax. The following model syntax and call to `lavaan()` achieves this:

```

R> HS.model.mixed <- '# latent variables
+                   visual  =~ 1*x1 + x2 + x3
+                   textual =~ 1*x4 + x5 + x6
+                   speed   =~ 1*x7 + x8 + x9
+                   # factor covariances
+                   visual  ~~ textual + speed
+                   textual  ~~ speed'
R> fit <- lavaan(HS.model.mixed, data = HolzingerSwineford1939,
+   auto.var = TRUE)

```

4.2. Examining the results

All three methods described above fit the same model. The `cfa()`, `sem()` and `lavaan()` fitting functions all return an object of class “lavaan”, for which several methods are available to examine model fit statistics and parameters estimates. Table 4 contains an overview of some of these methods.

The summary() method

Perhaps the most useful method to view results from a SEM fitted with **lavaan** is `summary()`. The `summary()` method can be called without any extra arguments, in which case only a short description of the model fit is displayed, together with the parameter estimates. Some extra arguments of the `summary()` method are `fit.measures`, `standardized`, and `rsquare`.

Method	Description
<code>summary()</code>	print a long summary of the model results
<code>show()</code>	print a short summary of the model results
<code>coef()</code>	returns the estimates of the free parameters in the model as a named numeric vector
<code>fitted()</code>	returns the implied moments (covariance matrix and mean vector) of the model
<code>resid()</code>	returns the raw, normalized or standardized residuals (difference between implied and observed moments)
<code>vcov()</code>	returns the covariance matrix of the estimated parameters
<code>predict()</code>	compute factor scores
<code>logLik()</code>	returns the log-likelihood of the fitted model (if maximum likelihood estimation was used)
<code>AIC()</code> , <code>BIC()</code>	compute information criteria (if maximum likelihood estimation was used)
<code>update()</code>	update a fitted lavaan object
<code>inspect()</code>	peek into the internal representation of the model; by default, it returns a list of model matrices counting the free parameters in the model; can also be used to extract starting values, gradient values, and much more

Table 4: Some methods for objects of class “lavaan”. See the help page for the **lavaan** class for more details (type `class?lavaan` at the R prompt).

If one or more of these is set to `TRUE`, the output will be enriched with additional fit measures, standardized estimates, and R^2 values for the dependent variables, respectively. In the example below, we request only the additional fit measures.

```
R> HS.model <- 'visual  =~ x1 + x2 + x3
+               textual =~ x4 + x5 + x6
+               speed   =~ x7 + x8 + x9'
R> fit <- cfa(HS.model, data = HolzingerSwineford1939)
R> summary(fit, fit.measures = TRUE)
```

lavaan (0.4-14) converged normally after 41 iterations

Number of observations	301
Estimator	ML
Minimum Function Chi-square	85.306
Degrees of freedom	24
P-value	0.000

Chi-square test baseline model:

Minimum Function Chi-square	918.852
Degrees of freedom	36
P-value	0.000

Full model versus baseline model:

Comparative Fit Index (CFI)	0.931
Tucker-Lewis Index (TLI)	0.896

Loglikelihood and Information Criteria:

Loglikelihood user model (H0)	-3737.745
Loglikelihood unrestricted model (H1)	-3695.092
Number of free parameters	21
Akaike (AIC)	7517.490
Bayesian (BIC)	7595.339
Sample-size adjusted Bayesian (BIC)	7528.739

Root Mean Square Error of Approximation:

RMSEA	0.092
90 Percent Confidence Interval	0.071 0.114
P-value RMSEA <= 0.05	0.001

Standardized Root Mean Square Residual:

```

SRMR                                0.065

Parameter estimates:

Information                                Expected
Standard Errors                          Standard

      Estimate  Std.err  Z-value  P(>|z|)
Latent variables:
visual =~
  x1          1.000
  x2          0.553    0.100    5.554    0.000
  x3          0.729    0.109    6.685    0.000
textual =~
  x4          1.000
  x5          1.113    0.065   17.014    0.000
  x6          0.926    0.055   16.703    0.000
speed =~
  x7          1.000
  x8          1.180    0.165    7.152    0.000
  x9          1.082    0.151    7.155    0.000

Covariances:
visual ~~
  textual      0.408    0.074    5.552    0.000
  speed        0.262    0.056    4.660    0.000
textual ~~
  speed        0.173    0.049    3.518    0.000

Variances:
  x1          0.549    0.114
  x2          1.134    0.102
  x3          0.844    0.091
  x4          0.371    0.048
  x5          0.446    0.058
  x6          0.356    0.043
  x7          0.799    0.081
  x8          0.488    0.074
  x9          0.566    0.071
visual       0.809    0.145
textual      0.979    0.112
speed       0.384    0.086

```

The output consists of three sections. The first section (the first 6 lines) contains the package version number, an indication whether the model has converged (and in how many iterations), and the effective number of observations used in the analysis. Next, the model χ^2 test statistic,

degrees of freedom, and a p value are printed. If `fit.measures = TRUE`, a second section is printed containing the test statistic of the baseline model (where all observed variables are assumed to be uncorrelated) and several popular fit indices. If maximum likelihood estimation is used, this section will also contain information about the loglikelihood, the AIC, and the BIC. The third section provides an overview of the parameter estimates, including the type of standard errors used and whether the observed or expected information matrix was used to compute standard errors. Then, for each model parameter, the estimate and the standard error are displayed, and if appropriate, a z value based on the Wald test and a corresponding two-sided p value are also shown. To ease the reading of the parameter estimates, they are grouped into three blocks: (1) factor loadings, (2) factor covariances, and (3) (residual) variances of both observed variables and factors.

The parameterEstimates() method

Although the `summary()` method provides a nice summary of the model results, it is useful for display only. An alternative is the `parameterEstimates()` method, which returns the parameter estimates as a `data.frame`, making the information easily accessible for further processing. By default, the `parameterEstimates()` method includes the estimates, standard errors, z value, p value, and 95% confidence intervals for all the model parameters.

```
R> parameterEstimates(fit)
```

	lhs	op	rhs	est	se	z	pvalue	ci.lower	ci.upper
1	visual	=~	x1	1.000	0.000	NA	NA	1.000	1.000
2	visual	=~	x2	0.553	0.100	5.554	0	0.358	0.749
3	visual	=~	x3	0.729	0.109	6.685	0	0.516	0.943
4	textual	=~	x4	1.000	0.000	NA	NA	1.000	1.000
5	textual	=~	x5	1.113	0.065	17.014	0	0.985	1.241
6	textual	=~	x6	0.926	0.055	16.703	0	0.817	1.035
7	speed	=~	x7	1.000	0.000	NA	NA	1.000	1.000
8	speed	=~	x8	1.180	0.165	7.152	0	0.857	1.503
9	speed	=~	x9	1.082	0.151	7.155	0	0.785	1.378
10	x1	~~	x1	0.549	0.114	4.833	0	0.326	0.772
11	x2	~~	x2	1.134	0.102	11.146	0	0.934	1.333
12	x3	~~	x3	0.844	0.091	9.317	0	0.667	1.022
13	x4	~~	x4	0.371	0.048	7.779	0	0.278	0.465
14	x5	~~	x5	0.446	0.058	7.642	0	0.332	0.561
15	x6	~~	x6	0.356	0.043	8.277	0	0.272	0.441
16	x7	~~	x7	0.799	0.081	9.823	0	0.640	0.959
17	x8	~~	x8	0.488	0.074	6.573	0	0.342	0.633
18	x9	~~	x9	0.566	0.071	8.003	0	0.427	0.705
19	visual	~~	visual	0.809	0.145	5.564	0	0.524	1.094
20	textual	~~	textual	0.979	0.112	8.737	0	0.760	1.199
21	speed	~~	speed	0.384	0.086	4.451	0	0.215	0.553
22	visual	~~	textual	0.408	0.074	5.552	0	0.264	0.552
23	visual	~~	speed	0.262	0.056	4.660	0	0.152	0.373
24	textual	~~	speed	0.173	0.049	3.518	0	0.077	0.270

The confidence level can be changed by setting the `level` argument. Setting `ci = FALSE` suppresses the confidence intervals. Another use of this function is to obtain several standardized versions of the estimates, by setting `standardized = TRUE`:

```
R> Est <- parameterEstimates(fit, ci = FALSE, standardized = TRUE)
R> subset(Est, op == "~")
```

	lhs	op	rhs	est	se	z	pvalue	std.lv	std.all	std.nox
1	visual	=~	x1	1.000	0.000	NA	NA	0.900	0.772	0.772
2	visual	=~	x2	0.553	0.100	5.554	0	0.498	0.424	0.424
3	visual	=~	x3	0.729	0.109	6.685	0	0.656	0.581	0.581
4	textual	=~	x4	1.000	0.000	NA	NA	0.990	0.852	0.852
5	textual	=~	x5	1.113	0.065	17.014	0	1.102	0.855	0.855
6	textual	=~	x6	0.926	0.055	16.703	0	0.917	0.838	0.838
7	speed	=~	x7	1.000	0.000	NA	NA	0.619	0.570	0.570
8	speed	=~	x8	1.180	0.165	7.152	0	0.731	0.723	0.723
9	speed	=~	x9	1.082	0.151	7.155	0	0.670	0.665	0.665

Here, only the factor loadings are shown. Relative to the prior output, three columns with standardized values were added. In the first column (`std.lv`), only the latent variables have been standardized; in the second column (`std.all`), both the latent and the observed variables have been standardized; in the third column (`std.nox`), both the latent and the observed variables have been standardized, except for the exogenous observed variables. The last of these options may be useful if the standardization of exogenous observed variables has little meaning (for example, binary covariates). Since there are no exogenous covariates in this model, the last two columns are identical in this output.

The `modificationIndices()` method

If the model fit is not superb, it may be informative to inspect the modification indices (MIs) and their corresponding expected parameter changes (EPCs). In essence, modification indices provide a rough estimate of how the χ^2 test statistic of a model would improve, if a particular parameter were unconstrained. The expected parameter change is the value this parameter would have if it were included as a free parameter. The `modificationIndices()` method (or the alias with the shorter name, `modindices()`) will print out a long list of parameters as a `data.frame`. In the output below, we only show those parameters for which the modification index is 10 or higher.

```
R> MI <- modificationIndices(fit)
R> subset(MI, mi > 10)
```

	lhs	op	rhs	mi	epc	sepc.lv	sepc.all	sepc.nox
1	visual	=~	x7	18.631	-0.422	-0.380	-0.349	-0.349
2	visual	=~	x9	36.411	0.577	0.519	0.515	0.515
3	x7	~~	x8	34.145	0.536	0.536	0.488	0.488
4	x8	~~	x9	14.946	-0.423	-0.423	-0.415	-0.415

The last three columns contain the standardized EPCs, using the same standardization conventions as for ordinary parameter estimates.




Figure 2: Path diagram of the structural equation model used to fit the Political Democracy data.

5. A second example: Structural equation modeling

In our second example, we will explore the ‘Industrialization and Political Democracy’ dataset previously used by Bollen in his 1989 book on structural equation modeling (Bollen 1989), and included with **lavaan** in the `PoliticalDemocracy` `data.frame`. The dataset contains various measures of political democracy and industrialization in developing countries. In the model, three latent variables are defined. The focus of the analysis is on the structural part of the model (i.e., the regressions among the latent variables). A graphical representation of the model is depicted in Figure 2.

5.1. Specifying and fitting the model, examining the results

For this example, we will only use the user-friendly `sem()` function to keep the model syntax as concise as possible. To describe the model, we need three different formula types: latent variables, regression formulas, and (co)variance formulas for the residual covariances among the observed variables. After the model has been fitted, we request a summary with no fit measures, but with standardized parameter estimates.

```

R> model <- '
+   # measurement model
+   ind60 =~ x1 + x2 + x3
+   dem60 =~ y1 + y2 + y3 + y4
+   dem65 =~ y5 + y6 + y7 + y8
+   # regressions
+   dem60 ~ ind60
+   dem65 ~ ind60 + dem60
+   # residual covariances
+   y1 ~~ y5
+   y2 ~~ y4 + y6
+   y3 ~~ y7

```

```

+      y4 ~~ y8
+      y6 ~~ y8'
R> fit <- sem(model, data = PoliticalDemocracy)
R> summary(fit, standardized = TRUE)

```

lavaan (0.4-14) converged normally after 70 iterations

Number of observations	75
Estimator	ML
Minimum Function Chi-square	38.125
Degrees of freedom	35
P-value	0.329

Parameter estimates:

	Estimate	Std.err	Z-value	P(> z)	Std.lv	Std.all
Latent variables:						
ind60 =~						
x1	1.000				0.670	0.920
x2	2.180	0.139	15.742	0.000	1.460	0.973
x3	1.819	0.152	11.967	0.000	1.218	0.872
dem60 =~						
y1	1.000				2.223	0.850
y2	1.257	0.182	6.889	0.000	2.794	0.717
y3	1.058	0.151	6.987	0.000	2.351	0.722
y4	1.265	0.145	8.722	0.000	2.812	0.846
dem65 =~						
y5	1.000				2.103	0.808
y6	1.186	0.169	7.024	0.000	2.493	0.746
y7	1.280	0.160	8.002	0.000	2.691	0.824
y8	1.266	0.158	8.007	0.000	2.662	0.828
Regressions:						
dem60 ~						
ind60	1.483	0.399	3.715	0.000	0.447	0.447
dem65 ~						
ind60	0.572	0.221	2.586	0.010	0.182	0.182
dem60	0.837	0.098	8.514	0.000	0.885	0.885
Covariances:						
y1 ~~						
y5	0.624	0.358	1.741	0.082	0.624	0.296
y2 ~~						

y4	1.313	0.702	1.871	0.061	1.313	0.273
y6	2.153	0.734	2.934	0.003	2.153	0.356
y3 ~~						
y7	0.795	0.608	1.308	0.191	0.795	0.191
y4 ~~						
y8	0.348	0.442	0.787	0.431	0.348	0.109
y6 ~~						
y8	1.356	0.568	2.386	0.017	1.356	0.338
Variances:						
x1	0.082	0.019			0.082	0.154
x2	0.120	0.070			0.120	0.053
x3	0.467	0.090			0.467	0.239
y1	1.891	0.444			1.891	0.277
y2	7.373	1.374			7.373	0.486
y3	5.067	0.952			5.067	0.478
y4	3.148	0.739			3.148	0.285
y5	2.351	0.480			2.351	0.347
y6	4.954	0.914			4.954	0.443
y7	3.431	0.713			3.431	0.322
y8	3.254	0.695			3.254	0.315
ind60	0.448	0.087			1.000	1.000
dem60	3.956	0.921			0.800	0.800
dem65	0.172	0.215			0.039	0.039

5.2. Parameter labels and simple equality constraints

In *lavaan*, each parameter has a name, referred to as the ‘parameter label’. The naming scheme is automatic and follows a simple set of rules. Each label consists of three components that describe the relevant formula defining the parameter. The first part is the variable name that appears on the left-hand side of the formula operator. The second part is the operator type of the formula, and the third part is the variable on the right-hand side of the operator that corresponds with the parameter. To see the naming mechanism in action, we can use the `coef()` function, which returns the (estimated) values of the free parameters and their corresponding parameter labels.

```
R> coef(fit)
```

ind60=~x2	ind60=~x3	dem60=~y2	dem60=~y3	dem60=~y4	dem65=~y6
2.180	1.819	1.257	1.058	1.265	1.186
dem65=~y7	dem65=~y8	dem60~ind60	dem65~ind60	dem65~dem60	y1~~y5
1.280	1.266	1.483	0.572	0.837	0.624
y2~~y4	y2~~y6	y3~~y7	y4~~y8	y6~~y8	x1~~x1
1.313	2.153	0.795	0.348	1.356	0.082
x2~~x2	x3~~x3	y1~~y1	y2~~y2	y3~~y3	y4~~y4
0.120	0.467	1.891	7.373	5.067	3.148

y5~~y5	y6~~y6	y7~~y7	y8~~y8	ind60~~ind60	dem60~~dem60
2.351	4.954	3.431	3.254	0.448	3.956
dem65~~dem65					
0.172					

Custom labels may be provided by the user in the model syntax, by pre-multiplying a variable name with that label. Consider, for example, the following regression formula:

```
y ~ b1*x1 + b2*x2 + b3*x3 + b4*x4
```

Here we have ‘named’ or ‘labeled’ the four regression coefficients as **b1**, **b2**, **b3** and **b4**. Custom labels are convenient because you can refer to them in other places of the model syntax. In particular, labels can be used to impose equality constraints on certain parameters. If two parameters have the same name, they will be considered to be the same, and only one value will be computed for them (i.e., a simple equality constraint). To illustrate this, we will respecify the model syntax of the Political Democracy data. In the original example in Bollen’s book, the factor loadings of the **dem60** factor are constrained to be equal to the factor loadings of the **dem65** factor. This make sense, since it is the same construct that is measured on two time points. To enforce these equality constraints, we label the factor loadings of the **dem60** factor (arbitrarily) as **d1**, **d2**, and **d3**. Note that we do not label the first factor loading since it is a fixed parameter (equal to 1.0). Next, we use the same labels for the factor loadings of the **dem65** factor, effectively imposing three equality constraints.

```
R> model.equal <- '# measurement model
+               ind60 =~ x1 + x2 + x3
+               dem60 =~ y1 + d1*y2 + d2*y3 + d3*y4
+               dem65 =~ y5 + d1*y6 + d2*y7 + d3*y8
+               # regressions
+               dem60 ~ ind60
+               dem65 ~ ind60 + dem60
+               # residual covariances
+               y1 ~~ y5
+               y2 ~~ y4 + y6
+               y3 ~~ y7
+               y4 ~~ y8
+               y6 ~~ y8'
R> fit.equal <- sem(model.equal, data = PoliticalDemocracy)
R> summary(fit.equal)
```

lavaan (0.4-14) converged normally after 69 iterations

Number of observations	75
Estimator	ML
Minimum Function Chi-square	40.179
Degrees of freedom	38
P-value	0.374

Parameter estimates:

Information				Expected
Standard Errors				Standard
	Estimate	Std.err	Z-value	P(> z)
Latent variables:				
ind60 =~				
x1	1.000			
x2	2.180	0.138	15.751	0.000
x3	1.818	0.152	11.971	0.000
dem60 =~				
y1	1.000			
y2 (d1)	1.191	0.139	8.551	0.000
y3 (d2)	1.175	0.120	9.755	0.000
y4 (d3)	1.251	0.117	10.712	0.000
dem65 =~				
y5	1.000			
y6 (d1)	1.191	0.139	8.551	0.000
y7 (d2)	1.175	0.120	9.755	0.000
y8 (d3)	1.251	0.117	10.712	0.000
Regressions:				
dem60 ~				
ind60	1.471	0.392	3.750	0.000
dem65 ~				
ind60	0.600	0.226	2.661	0.008
dem60	0.865	0.075	11.554	0.000
Covariances:				
y1 ~~				
y5	0.583	0.356	1.637	0.102
y2 ~~				
y4	1.440	0.689	2.092	0.036
y6	2.183	0.737	2.960	0.003
y3 ~~				
y7	0.712	0.611	1.165	0.244
y4 ~~				
y8	0.363	0.444	0.817	0.414
y6 ~~				
y8	1.372	0.577	2.378	0.017
Variances:				
x1	0.081	0.019		
x2	0.120	0.070		
x3	0.467	0.090		

y1	1.855	0.433
y2	7.581	1.366
y3	4.956	0.956
y4	3.225	0.723
y5	2.313	0.479
y6	4.968	0.921
y7	3.560	0.710
y8	3.308	0.704
ind60	0.449	0.087
dem60	3.875	0.866
dem65	0.164	0.227

The fit of the constrained model is slightly worse compared to the unconstrained model. But is it significantly worse? To compare two nested models, we can use the `anova()` function, which will compute the χ^2 difference test:

```
R> anova(fit, fit.equal)
```

Chi Square Difference Test

	Df	AIC	BIC	Chisq	Chisq diff	Df diff	Pr(>Chisq)
fit	35	3157.6	3229.4	38.125			
fit.equal	38	3153.6	3218.5	40.179	2.0543	3	0.5612

5.3. Extracting fit measures

The `summary()` method with the argument `fit.measures = TRUE` will output a number of fit measures. If fit statistics are needed for further processing, the `fitMeasures()` method is preferred. The first argument to `fitMeasures()` is the fitted object and the second argument is a character vector containing the names of the fit measures one wish to extract. For example, if we only need the CFI and RMSEA values, we can use:

```
R> fitMeasures(fit, c("cfi", "rmsea"))
```

```
  cfi rmsea
0.995 0.035
```

Omitting the second argument will return all the fit measures computed by **lavaan**.

5.4. Using the `inspect()` method

To finish our SEM example, we will briefly mention the `inspect()` method which allows the user to peek under the hood of a **lavaan** object. By default, calling `inspect()` on a fitted **lavaan** object returns a list of the model matrices that are used internally to represent the model. The free parameters are nonzero integers.

```
R> inspect(fit)
```

```

$lambda
      ind60 dem60 dem65
x1      0      0      0
x2      1      0      0
x3      2      0      0
y1      0      0      0
y2      0      3      0
y3      0      4      0
y4      0      5      0
y5      0      0      0
y6      0      0      6
y7      0      0      7
y8      0      0      8

$theta
      x1 x2 x3 y1 y2 y3 y4 y5 y6 y7 y8
x1 18
x2 0 19
x3 0 0 20
y1 0 0 0 21
y2 0 0 0 0 22
y3 0 0 0 0 0 23
y4 0 0 0 0 13 0 24
y5 0 0 0 12 0 0 0 25
y6 0 0 0 0 14 0 0 0 26
y7 0 0 0 0 0 15 0 0 0 27
y8 0 0 0 0 0 0 16 0 17 0 28

$psi
      ind60 dem60 dem65
ind60 29
dem60 0 30
dem65 0 0 31

$beta
      ind60 dem60 dem65
ind60 0 0 0
dem60 9 0 0
dem65 10 11 0

```

The output reveals that **lavaan** is currently using the **LISREL** matrix representation, albeit with no distinction between endogenous and exogenous variables. This is the so-called ‘all-y’ representation. In future releases, I plan to allow for alternative matrix representations, including the Bentler-Weeks and the reticular action model (RAM) approach ([Bollen 1989](#), chapter 9). To see the starting values of the parameters in each model matrix, type

```
R> inspect(fit, what = "start")
```



```

$lambda
      ind60 dem60 dem65
x1 1.000 0.000 0.000
x2 2.193 0.000 0.000
x3 1.824 0.000 0.000
y1 0.000 1.000 0.000
y2 0.000 1.296 0.000
y3 0.000 1.055 0.000
y4 0.000 1.294 0.000
y5 0.000 0.000 1.000
y6 0.000 0.000 1.303
y7 0.000 0.000 1.403
y8 0.000 0.000 1.401

$theta
      x1      x2      x3      y1      y2      y3      y4      y5      y6      y7      y8
x1 0.265
x2 0.000 1.126
x3 0.000 0.000 0.975
y1 0.000 0.000 0.000 3.393
y2 0.000 0.000 0.000 0.000 7.686
y3 0.000 0.000 0.000 0.000 0.000 5.310
y4 0.000 0.000 0.000 0.000 0.000 0.000 5.535
y5 0.000 0.000 0.000 0.000 0.000 0.000 0.000 3.367
y6 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 5.612
y7 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 5.328
y8 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 5.197

$psi
      ind60 dem60 dem65
ind60 0.05
dem60 0.00 0.05
dem65 0.00 0.00 0.05

$beta
      ind60 dem60 dem65
ind60      0      0      0
dem60      0      0      0
dem65      0      0      0

```

Many more inspect options are described in the help page for the `lavaan` class.

6. Multiple groups

The **lavaan** package has full support for multiple-groups SEM. To request a multiple-groups analysis, the variable defining group membership in the dataset can be passed to the `group`

argument of the `cfa()`, `sem()`, or `lavaan()` function calls. By default, the same model is fitted in all groups without any equality constraints on the model parameters. In the following example, we fit the H&S model for the two schools (Pasteur and Grant-White).

```
R> HS.model <- 'visual =~ x1 + x2 + x3
+               textual =~ x4 + x5 + x6
+               speed  =~ x7 + x8 + x9'
R> fit <- cfa(HS.model, data = HolzingerSwineford1939, group = "school")
```

The `summary()` output is rather long and not shown here. Essentially, it shows a set of parameter estimates for the Pasteur group, followed by another set of parameter estimates for the Grant-White group. If we wish to impose equality constraints on model parameters across groups, we can use the `group.equal` argument. For example, `group.equal = c("loadings", "intercepts")` will constrain both the factor loadings and the intercepts of the observed variables to be equal across groups. Other options that can be included in the `group.equal` argument are described in the help pages of the fitting functions. As a short example, we will fit the H&S model for the two schools, but constrain the factor loadings and intercepts to be equal. The `anova` function can be used to compare the two model fits.

```
R> fit.metric <- cfa(HS.model, data = HolzingerSwineford1939,
+   group = "school", group.equal = c("loadings", "intercepts"))
R> anova(fit, fit.metric)
```

Chi Square Difference Test

	Df	AIC	BIC	Chisq	Chisq diff	Df diff	Pr(>Chisq)
fit	48	7484.4	7706.8	115.85			
fit.metric	60	7508.6	7686.6	164.10	48.251	12	2.826e-06

If the `group.equal` argument is used to constrain the factor loadings across groups, all factor loadings are affected. If some exceptions are needed, one can use the `group.partial` argument, which takes a vector of parameter labels that specify which parameters will remain free across groups. Therefore, the combination of the `group.equal` and `group.partial` arguments gives the user a flexible mechanism to specify across group equality constraints.

7. Can lavaan do this? A short feature list

The **lavaan** package has many features, and we foresee that the feature list will keep growing in the next couple of years. To present the reader a flavor of the current capabilities of **lavaan**, I will use this section to mention briefly a variety of topics that are often of interest to users of SEM software.

7.1. Non-normal and categorical data

The 0.4 version of the **lavaan** package only supports continuous data. Support for categorical data is expected in the 0.5 release. In the current release, however, I have devoted considerable attention to dealing with non-normal continuous data. In the SEM literature, the topic of

dealing with non-normal data is well studied (see [Finney and DiStefano 2006](#), for an overview). Three popular strategies to deal with non-normal data have been implemented in the **lavaan** package: asymptotically distribution-free estimation; scaled test statistics and robust standard errors; and bootstrapping.

Asymptotically distribution-free (ADF) estimation

The asymptotically distribution-free (ADF) estimator ([Browne 1984](#)) makes no assumption of normality. Therefore, variables that are skewed or kurtotic do not distort the ADF based standard errors and test statistic. The ADF estimator is part of a larger family of estimators called weighted least squares (WLS) estimators. The discrepancy function of these estimators can be written as follows:

$$F = (\mathbf{s} - \hat{\boldsymbol{\sigma}})^\top \mathbf{W}^{-1} (\mathbf{s} - \hat{\boldsymbol{\sigma}})$$

where \mathbf{s} is a vector of the non-duplicated elements in the sample covariance matrix (\mathbf{S}), $\hat{\boldsymbol{\sigma}}$ is a vector of the non-duplicated elements of the model-implied covariance matrix ($\hat{\boldsymbol{\Sigma}}$), and \mathbf{W} is a weight matrix. The weight matrix utilized with the ADF estimator is the asymptotic covariance matrix: a matrix of the covariances of the observed sample variances and covariances. Theoretically, the ADF estimator seems to be a perfect strategy to deal with non-normal data. Unfortunately, empirical research (e.g., [Olsson, Foss, Troye, and Howell 2000](#)) has shown that the ADF method breaks down unless the sample size is huge (e.g., $N > 5000$).

In **lavaan**, the estimator can be set by using the `estimator` argument in one of the fitting functions. The default is maximum likelihood estimation, or `estimator = "ML"`. To switch to the ADF estimator, you can set `estimator = "WLS"`.

Satorra-Bentler scaled test statistic and robust standard errors

An alternative strategy is to use maximum likelihood (ML) for estimating the model parameters, even if the data are known to be non-normal. In this case, the parameter estimates are still consistent (if the model is identified and correctly specified), but the standard errors tend to be too small (as much as 25–50%), meaning that we may reject the null hypothesis (that a parameter is zero) too often. In addition, the model (χ^2) test statistic tends to be too large, meaning that we may reject the model too often.

In the SEM literature, several authors have extended the ML methodology to produce standard errors that are asymptotically correct for arbitrary distributions (with finite fourth-order moments), and where a rescaled test statistic is used for overall model evaluation.

Standard errors of the maximum likelihood estimators are based on the covariance matrix that is obtained by inverting the associated information matrix. Robust standard errors replace this covariance matrix by a sandwich-type covariance matrix, first proposed by [Huber \(1967\)](#) and introduced in the SEM literature by [Bentler \(1983\)](#); [Browne \(1984\)](#); [Browne and Arminger \(1995\)](#) and many others. In **lavaan**, the `se` argument can be used to switch between different types of standard errors. Setting `se = "robust"` will produce robust standard errors based on a sandwich-type covariance matrix.

The best known method for correcting the model test statistic is the so-called Satorra-Bentler scaled test statistic ([Satorra and Bentler 1988, 1994](#)). Their method rescales the value of the ML-based χ^2 test statistic by an amount that reflects the degree of kurtosis. Several

simulation studies have shown that this correction is effective with non-normal data (Chou, Bentler, and Satorra 1991; Curran, West, and Finch 1996), even in small to moderate samples. And because applying the correction often results in a better model fit, it is not surprising that the Satorra-Bentler rescaling method has become popular among SEM users.

In **lavaan**, the `test` argument can be used to switch between different test statistics. Setting `test = "satorra.bentler"` supplements the standard χ^2 model test with the scaled version. In the output produced by the `summary()` method, both scaled and unscaled model tests (and corresponding fit indices) are displayed in adjacent columns. Because one typically wants both robust standard errors and a scaled test statistic, specifying `estimator = "MLM"` fits the model using standard maximum likelihood to estimate the model parameters, but with robust standard errors and a Satorra-Bentler scaled test statistic.

```
R> fit <- cfa(HS.model, data = HolzingerSwineford1939, missing = "listwise",
+   estimator = "MLM", mimic = "Mplus")
R> summary(fit, estimates = FALSE, fit.measures = TRUE)
```

lavaan (0.4-14) converged normally after 41 iterations

Number of observations	301	
Estimator	ML	Robust
Minimum Function Chi-square	85.306	81.908
Degrees of freedom	24	24
P-value	0.000	0.000
Scaling correction factor		1.041
for the Satorra-Bentler correction (Mplus variant)		

Chi-square test baseline model:

Minimum Function Chi-square	918.852	888.912
Degrees of freedom	36	36
P-value	0.000	0.000

Full model versus baseline model:

Comparative Fit Index (CFI)	0.931	0.932
Tucker-Lewis Index (TLI)	0.896	0.898

Loglikelihood and Information Criteria:

Loglikelihood user model (H0)	-3737.745	-3737.745
Loglikelihood unrestricted model (H1)	-3695.092	-3695.092
Number of free parameters	30	30
Akaike (AIC)	7535.490	7535.490
Bayesian (BIC)	7646.703	7646.703
Sample-size adjusted Bayesian (BIC)	7551.560	7551.560

Root Mean Square Error of Approximation:

RMSEA		0.092	0.090	
90 Percent Confidence Interval	0.071	0.114	0.069	0.111
P-value RMSEA <= 0.05		0.001	0.001	

Standardized Root Mean Square Residual:

SRMR	0.060	0.060
------	-------	-------

In this example, the `mimic = "Mplus"` argument was used to mimic the way the **Mplus** program computes the Satorra-Bentler scaled test statistic. By default (i.e., when the `mimic` argument is omitted), **lavaan** will use the method that is used by the **EQS** program. To mimic the exact value of the Satorra-Bentler scaled test statistic as reported by the **EQS** program, one can use

```
R> fit <- cfa(HS.model, data = HolzingerSwineford1939, estimator = "MLM",
+   mimic = "EQS")
R> fit
```

lavaan (0.4-14) converged normally after 41 iterations

Number of observations	301	
Estimator	ML	Robust
Minimum Function Chi-square	85.022	81.141
Degrees of freedom	24	24
P-value	0.000	0.000
Scaling correction factor		1.048
for the Satorra-Bentler correction		

If two models are nested, but the model test statistics are scaled, the usual χ^2 difference test can no longer be used. Instead, a special procedure is needed known as the scaled χ^2 difference test (Satorra and Bentler 2001). The `anova()` function in **lavaan** will automatically detect this and compute a scaled χ^2 difference test if appropriate.

Bootstrapping: The naïve bootstrap and the Bollen-Stine bootstrap

The third strategy for dealing with non-normal data is bootstrapping. For standard errors, we can use the standard nonparametric bootstrap to obtain bootstrap standard errors. However, to bootstrap the test statistic (and its *p* value), the standard (naïve) bootstrap is incorrect because it reflects not only non-normality and sampling variability, but also model misfit. Therefore, the original sample must first be transformed so that the sample covariance matrix corresponds with the model-implied covariance. In the SEM literature, this model-based bootstrap procedure is known as the Bollen-Stine bootstrap (Bollen and Stine 1993).

In **lavaan**, bootstrap standard errors can be obtained by setting `se = "bootstrap"`. In this case, the confidence intervals produced by the `parameterEstimates()` method will be

bootstrap-based confidence intervals. If `test = "bootstrap"` or `test = "bollen.stine"`, the data are first transformed to perform a model-based ‘Bollen-Stine’ bootstrap. The bootstrap standard errors are also based on these model-based bootstrap draws, and the standard p value of the χ^2 test statistic is supplemented with a bootstrap probability value, obtained by computing the proportion of test statistics from the bootstrap samples that exceed the value of the test statistic from the original (parent) sample.

By default, **lavaan** generates $R = 1000$ bootstrap draws, but this number can be changed by setting the `bootstrap` argument. It may be informative to set `verbose = TRUE` to monitor the progress of bootstrapping.

7.2. Missing data

If the data contain missing values, the default behavior in **lavaan** is listwise deletion. If the missing mechanism is MCAR (missing completely at random) or MAR (missing at random), the **lavaan** package provides case-wise (or ‘full information’) maximum likelihood (FIML) estimation (Arbuckle 1996). FIML estimation can be enabled by specifying the argument `missing = "ml"` (or its alias `missing = "fiml"`) when calling the fitting function. An unrestricted (h1) model will automatically be estimated, so that all common fit indices are available. Robust standard errors are also available, as is a scaled test statistic (Yuan and Bentler 2000) if the data are both incomplete and non-normal.

7.3. Linear and nonlinear equality and inequality constraints

In many applications, it is necessary to impose constraints on some of the model parameters. For example, one may wish to enforce that a variance parameter is strictly positive. For some models, it is important to specify that a parameter is equal to some (linear or nonlinear) function of other parameters. The aim of the **lavaan** package is to make such constraints easy to specify using the **lavaan** model syntax. A short example will illustrate constraint syntax in **lavaan**. Consider the following regression:

$$y \sim b_1 \cdot x_1 + b_2 \cdot x_2 + b_3 \cdot x_3$$

where we have explicitly labeled the regression coefficients as b_1 , b_2 and b_3 . We create a toy dataset containing these four variables and fit the regression model:

```
R> set.seed(1234)
R> Data <- data.frame(y = rnorm(100), x1 = rnorm(100), x2 = rnorm(100),
+   x3 = rnorm(100))
R> model <- 'y ~ b1*x1 + b2*x2 + b3*x3'
R> fit <- sem(model, data = Data)
R> coef(fit)
```

	b_1	b_2	b_3	$y \sim y$
	-0.052	0.084	0.139	0.970

Suppose that we wish to impose two (nonlinear) constraints on b_1 : $b_1 = (b_2 + b_3)^2$ and $b_1 \geq \exp(b_2 + b_3)$. The first constraint is an equality constraint, whereas the second is an inequality constraint. Both constraints are nonlinear. In **lavaan**, this is accomplished as follows:

```
lavaan (0.4-14) converged normally after 49 iterations
```

Parameter estimates:

		Estimate	Std.err	Z-value	P(> z)
Regressions:					
y ~					
x1	(b1)	0.495			
x2	(b2)	-0.405	0.092	-4.411	0.000
x3	(b3)	-0.299	0.092	-3.256	0.001

```

Variances:
      y      1.610      0.228

```

```
Constraints:                                     Slack (>=0)
      b1 - (exp(b2+b3))                        0.000
      b1 - ((b2+b3)^2)                        0.000
```

7.4. Indirect effects and mediation analysis

Once a model has been fitted, we may be interested in values that are functions of the original estimated model parameters. One example is an indirect effect which is a product of two (or more) regression coefficients. Consider a classical mediation setup with three variables: \mathbf{Y} is the dependent variable, \mathbf{X} is the predictor, and \mathbf{M} is a mediator. For illustration, we again create

a toy dataset containing these three variables, and fit a path analysis model that includes the direct effect of X on Y and the indirect effect of X on Y via M.

```
R> set.seed(1234)
R> X <- rnorm(100)
R> M <- 0.5 * X + rnorm(100)
R> Y <- 0.7 * M + rnorm(100)
R> Data <- data.frame(X = X, Y = Y, M = M)
R> model <- '# direct effect
+           Y ~ c*X
+           # mediator
+           M ~ a*X
+           Y ~ b*M
+           # indirect effect (a*b)
+           ab := a*b
+           # total effect
+           total := c + (a*b)'
R> fit <- sem(model, data = Data)
R> summary(fit)
```

lavaan (0.4-14) converged normally after 13 iterations

Number of observations	100
Estimator	ML
Minimum Function Chi-square	0.000
Degrees of freedom	0
P-value	0.000

Parameter estimates:

	Information		Expected		
	Standard Errors		Standard		
	Estimate	Std.err	Z-value	P(> z)	
Regressions:					
Y ~					
X	(c)	0.036	0.104	0.348	0.728
M ~					
X	(a)	0.474	0.103	4.613	0.000
Y ~					
M	(b)	0.788	0.092	8.539	0.000
Variances:					
Y		0.898	0.127		
M		1.054	0.149		

Defined parameters:

ab	0.374	0.092	4.059	0.000
total	0.410	0.125	3.287	0.001

The example illustrates the use of the ‘:=’ operator in the **lavaan** model syntax. This operator ‘defines’ new parameters which take on values that are an arbitrary function of the original model parameters. The function, however, must be specified in terms of the parameter *labels* that are explicitly mentioned in the model syntax. By default, the standard errors for these defined parameters are computed using the delta method (Sobel 1982). As with other models, bootstrap standard errors can be requested simply by specifying `se = "bootstrap"` in the fitting function.

8. Concluding remarks

This paper described the R package **lavaan**. Despite its name, the current version (0.4) of **lavaan** should be regarded as a package for structural equation modeling with continuous data. One of the main attractions of **lavaan** is its intuitive and easy-to-use model syntax. **lavaan** is also fairly complete, and contains most of the features that applied researchers are looking for in a modern SEM package. So when will **lavaan** become a package for latent variable analysis? In due time.

References

- Anderson TW, Rubin H (1956). “Statistical Inference in Factor Analysis.” In *Proceedings of the Third Berkeley Symposium on Mathematical Statistics and Probability*, pp. 111–150. University of California Press, Berkeley.
- Arbuckle JL (1996). “Full Information Estimation in the Presence of Incomplete Data.” In GA Marcoulides, RE Schumacker (eds.), *Advanced Structural Equation Modeling*, pp. 243–277. Lawrence Erlbaum, Mahwah.
- Arbuckle JL (2011). *IBM SPSS AMOS 20 User’s Guide*. IBM Corporation, Armonk.
- Bentler PM (1983). “Some Contributions to Efficient Statistics in Structural Models: Specification and Estimation of Moment Structures.” *Psychometrika*, **48**, 493–517.
- Bentler PM (2004). *EQS 6 Structural Equations Program Book*. Multivariate Software, Inc., Encino.
- Boker S, Neale M, Maes HH, Wilde M, Spiegel M, Brick T, Spies J, Estabrook R, Kenny S, Bates T, Mehta P, Fox J (2011). “**OpenMx**: An Open Source Extended Structural Equation Modeling Framework.” *Psychometrika*, **76**, 306–317.
- Bollen KA (1989). *Structural Equations with Latent Variables*. John Wiley & Sons.
- Bollen KA, Stine RA (1993). “Bootstrapping Goodness-of-fit Measures in Structural Equation Models.” In KA Bollen, JS Long (eds.), *Testing Structural Equation Models*, pp. 111–135. Sage Publications, Newbury Park.

- Browne MW (1974). “Generalized Least Squares Estimators in the Analysis of Covariances Structures.” *South African Statistical Journal*, **8**, 1–24.
- Browne MW (1984). “Asymptotic Distribution-Free Methods in the Analysis of Covariance Structures.” *British Journal of Mathematical and Statistical Psychology*, **37**, 62–83.
- Browne MW, Arminger G (1995). “Specification and Estimation of Mean- and Covariance-Structure Models.” In G Arminger, CC Clogg, ME Sobel (eds.), *Handbook of Statistical Modeling for the Social and Behavioral Sciences*, pp. 311–359. Plenum Press, New York.
- Chou CP, Bentler PM, Satorra A (1991). “Scaled Test Statistics and Robust Standard Errors for Nonnormal Data in Covariance Structure Analysis: A Monte-Carlo Study.” *British Journal of Mathematical and Statistical Psychology*, **44**, 347–357.
- Curran PJ, West SG, Finch JF (1996). “The Robustness of Test Statistics to Non-Normality and Specification Error in Confirmatory Factor Analysis.” *Psychological Methods*, **1**, 16–29.
- Duncan OD (1966). “Path Analysis: Sociological Examples.” *American Journal of Sociology*, **72**, 1–16.
- Finney SJ, DiStefano C (2006). “Non-Normal and Categorical Data in Structural Equation Modeling.” In GR Hancock, RO Mueller (eds.), *Structural Equation Modeling: A Second Course*, pp. 269–314. Information Age Publishing.
- Fox J (2006). “Structural Equation Modeling with the **sem** Package in R.” *Structural Equation Modeling: A Multidisciplinary Journal*, **13**, 465–486.
- Fox J, Nie Z, Byrnes J (2012). *sem: Structural Equation Models*. R package version 3.0-0, URL <http://CRAN.R-project.org/package=sem>.
- Haavelmo T (1943). “The Statistical Implications of a System of Simultaneous Equations.” *Econometrica*, **11**, 1–12.
- Hallquist M (2012). *MplusAutomation: Automating Mplus Model Estimation and Interpretation*. R package version 0.5-1, URL <http://CRAN.R-project.org/package=MplusAutomation>.
- Hauser RM, Goldberger AS (1971). “The Treatment of Unobservable Variables in Path Analysis.” *Sociological Methodology*, **3**, 81–117.
- Huber PJ (1967). “The Behavior of Maximum Likelihood Estimates under Nonstandard Conditions.” In *Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability*, volume 1, pp. 221–233. University of California Press, Berkeley.
- Jöreskog KG (1970). “A General Method for Analysis of Covariance Structures.” *Biometrika*, **57**, 239–251.
- Jöreskog KG (1973). “A General Method for Estimating a Linear Structural Equation System.” In AS Goldberger, OD Duncan (eds.), *Structural Equation Models in the Social Sciences*, pp. 85–112. Seminar Press, New York.
- Jöreskog KG, Sörbom D (1997). *LISREL 8: User’s Reference Guide*. Scientific Software International.

- Keesling JW (1973). *Maximum Likelihood Approaches to Causal Flow Analysis*. Ph.D. thesis, Department of Education, University of Chicago.
- Koopmans T (1945). “Statistical Estimation of Simultaneous Economic Relations.” *Journal of the American Statistical Association*, **40**, 448–466.
- Lawley DN (1940). “The Estimation of Factor Loadings by the Method of Maximum Likelihood.” In *Proceedings of the Royal Society of Edinburgh*, volume 60, pp. 64–82.
- Lee SY (2007). *Handbook of Latent Variable and Related Models*. Elsevier, Amsterdam.
- Mair P, Wu E, Bentler PM (2010). “**EQS** Goes R: Embedding **EQS** into the R Environment Using the Package **REQS**.” *Structural Equation Modeling: A Multidisciplinary Journal*, **17**, 333–349.
- Muthén BO (2002). “Beyond SEM: General Latent Variable Modeling.” *Behaviormetrika*, **29**, 81–117.
- Muthén LK, Muthén BO (2010). *Mplus User’s Guide*. Muthén & Muthén, Los Angeles, 6th edition.
- Neale MC, Boker SM, Xie G, Maes HH (2003). *Mx: Statistical Modeling*. Department of Psychiatry, VCU Box 900126, Richmond, VA 23298, 6th edition.
- Olsson UH, Foss T, Troye S, Howell R (2000). “The Performance of ML, GLS, and WLS Estimation in Structural Equation Modeling under Conditions of Misspecification and Non-normality.” *Structural Equation Modeling: A Multidisciplinary Journal*, **7**, 557–595.
- Rabe-Hesketh S, Skrondal A, Pickles A (2004). “Generalized Multilevel Structural Equation Modelling.” *Psychometrika*, **69**, 167–190.
- R Development Core Team (2012). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. ISBN 3-900051-07-0, URL <http://www.R-project.org/>.
- Satorra A, Bentler PM (1988). “Scaling Corrections for Chi-Square Statistics in Covariance Structure Analysis.” In *ASA 1988 Proceedings of the Business and Economic Statistics Section*, volume 1, pp. 308–313. American Statistical Association, Alexandria.
- Satorra A, Bentler PM (1994). “Corrections to Test Statistics and Standard Errors in Covariance Structure Analysis.” In A von Eye, CC Clogg (eds.), *Latent Variables Analysis: Applications for Developmental Research*, pp. 399–419. Sage, Thousands Oaks.
- Satorra A, Bentler PM (2001). “A Scaled Difference Chi-Square Test Statistic for Moment Structure Analysis.” *Psychometrika*, **66**, 507–514.
- Skrondal A, Rabe-Hesketh S (2004). *Generalized Latent Variable Modeling: Multilevel, Longitudinal, and Structural Equation Models*. Chapman & Hall/CRC, Boca Raton.
- Sobel ME (1982). “Asymptotic Confidence Intervals for Indirect Effects in Structural Equation Models.” *Sociological Methodology*, **13**, 290–312.

- Spearman C (1904). “General Intelligence, Objectively Determined and Measured.” *The American Journal of Psychology*, **15**, 201–292.
- van de Schoot R, Hoijsink H, Deković M (2010). “Testing Inequality Constrained Hypotheses in SEM Models.” *Structural Equation Modeling: A Multidisciplinary Journal*, **17**, 443–463.
- Wiley DE (1973). “The Identification Problem for Structural Equation Models with Unmeasured Variables.” In AS Goldberger, OD Duncan (eds.), *Structural Equation Models in the Social Sciences*, pp. 69–84. Seminar Press, New York.
- Wright S (1921). “Correlation and Causation.” *Journal of Agricultural Research*, **20**, 557–585.
- Yuan KH, Bentler PM (2000). “Three Likelihood-Based Methods for Mean and Covariance Structure Analysis with Nonnormal Missing Data.” *Sociological Methodology*, **30**, 165–200.
- Zellner A (1970). “Estimation of Regression Relationships Containing Unobservable Independent Variables.” *International Economic Review*, **11**, 441–454.

Affiliation:

Yves Rosseel
Department of Data Analysis
Ghent University
9000 Ghent, Belgium
E-mail: Yves.Rosseel@UGent.be
URL: <http://lavaan.org/>